

```
In [1]: import pandas as pd
```

```
In [2]: import warnings
warnings.filterwarnings("ignore")
```

```
In [3]: data=pd.read_csv("/home/placement/Desktop/EEE(222)/fiat500.csv")
```

```
In [4]: data.describe()
```

```
Out[4]:
```

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	price
count	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000
mean	769.500000	51.904421	1650.980494	53396.011704	1.123537	43.541361	11.563428	8576.003901
std	444.126671	3.988023	1289.522278	40046.830723	0.416423	2.133518	2.328190	1939.958641
min	1.000000	51.000000	366.000000	1232.000000	1.000000	36.855839	7.245400	2500.000000
25%	385.250000	51.000000	670.000000	20006.250000	1.000000	41.802990	9.505090	7122.500000
50%	769.500000	51.000000	1035.000000	39031.000000	1.000000	44.394096	11.869260	9000.000000
75%	1153.750000	51.000000	2616.000000	79667.750000	1.000000	45.467960	12.769040	10000.000000
max	1538.000000	77.000000	4658.000000	235000.000000	4.000000	46.795612	18.365520	11100.000000

```
In [5]: data=data.drop(['ID','lat','lon'],axis=1)
```

```
In [6]: data
```

```
Out[6]:
```

	model	engine_power	age_in_days	km	previous_owners	price
0	lounge	51	882	25000	1	8900
1	pop	51	1186	32500	1	8800
2	sport	74	4658	142228	1	4200
3	lounge	51	2739	160000	1	6000
4	pop	73	3074	106880	1	5700
...
1533	sport	51	3712	115280	1	5200
1534	lounge	74	3835	112000	1	4600
1535	pop	51	2223	60457	1	7500
1536	lounge	51	2557	80750	1	5990
1537	pop	51	1766	54276	1	7900

1538 rows × 6 columns

```
In [7]: data=pd.get_dummies(data)
```

```
In [8]: data.shape
```

```
Out[8]: (1538, 8)
```

```
In [9]: data
```

```
Out[9]:
```

	engine_power	age_in_days	km	previous_owners	price	model_lounge	model_pop	model_sport
0	51	882	25000	1	8900	1	0	0
1	51	1186	32500	1	8800	0	1	0
2	74	4658	142228	1	4200	0	0	1
3	51	2739	160000	1	6000	1	0	0
4	73	3074	106880	1	5700	0	1	0
...
1533	51	3712	115280	1	5200	0	0	1
1534	74	3835	112000	1	4600	1	0	0
1535	51	2223	60457	1	7500	0	1	0
1536	51	2557	80750	1	5990	1	0	0
1537	51	1766	54276	1	7900	0	1	0

1538 rows × 8 columns

```
In [10]: y=data['km']  
x=data.drop('km',axis=1)
```

In [11]: y

```
Out[11]: 0      25000
          1      32500
          2     142228
          3     160000
          4     106880
          ...
        1533    115280
        1534    112000
        1535     60457
        1536     80750
        1537     54276
        Name: km, Length: 1538, dtype: int64
```

In [12]: x

```
Out[12]:
```

	engine_power	age_in_days	previous_owners	price	model_lounge	model_pop	model_sport
0	51	882	1	8900	1	0	0
1	51	1186	1	8800	0	1	0
2	74	4658	1	4200	0	0	1
3	51	2739	1	6000	1	0	0
4	73	3074	1	5700	0	1	0
...
1533	51	3712	1	5200	0	0	1
1534	74	3835	1	4600	1	0	0
1535	51	2223	1	7500	0	1	0
1536	51	2557	1	5990	1	0	0
1537	51	1766	1	7900	0	1	0

1538 rows × 7 columns

```
In [13]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33,random_state=42)
```

```
In [14]: x_test.head(5)
```

```
Out[14]:
```

	engine_power	age_in_days	previous_owners	price	model_lounge	model_pop	model_sport
481	51	3197	2	7900	0	1	0
76	62	2101	1	7900	0	1	0
1502	51	670	1	9400	1	0	0
669	51	913	1	8500	1	0	0
1409	51	762	1	9700	1	0	0

```
In [15]: y_train.head(10)
```

```
Out[15]: 527    13111
129     21400
602     57039
331     40700
323     16783
1358    29378
522     18443
584     11997
1236    66900
535     35000
Name: km, dtype: int64
```

linear regression

```
In [16]: from sklearn.linear_model import LinearRegression
reg=LinearRegression()#creating object of linearregression
reg.fit(x_train,y_train)#traning and fitting lr object using traning data
```

Out[16]: LinearRegression()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [17]: ypred=reg.predict(x_test)
```

```
In [18]: ypred
```

```
Out[18]: array([ 75731.24452341,  66029.84772251,  34401.2360973 ,  47789.21545234,
  31445.26338572,  28940.45224947,  29985.55310965,  27782.98986147,
  18389.94810051,  43647.33173542,  26400.6856128 ,  65315.51718461,
  78520.35323068,  79990.40592647,  40933.24832683,  17977.18600208,
  21520.86799789,  86390.77726774, 134696.69941688,  18247.17039096,
  28441.83272222,  17977.18600208,  84917.61592634,  31095.18096823,
 102601.42053493,  18683.76101173, 126929.83769068,  80736.28377642,
  61845.82202184,  25427.09241246,  75352.29888237, 107607.91154532,
 102849.71048838, 131742.45012204,  54385.26759139,  96392.07925672,
  17490.98815001,  72194.05616406, 102326.85312655,  58616.36626514,
  23192.26566508,  82637.99477227,  27133.01992404,  22420.51239391,
  61892.93067256,  12228.9994177 ,  44144.63447273,  33196.66505858,
 107132.46976595,  66916.67221792,  38015.67509214,  13863.75090251,
  18037.8995605 ,  83477.63824197,  40401.9887229 ,  30272.50416449,
  27657.65528172,  13593.76651363,  21795.11626488,  57818.16924686,
  21156.03445181,  92459.0954649 ,  25427.09241246,  83317.89386646,
  82265.53558981,  28670.46786059,  29579.53406254,  36591.80158707,
  36351.95918295, 103152.37732341,  82659.4275832 ,  96271.05349949,
  51009.76245379,  18247.17039096,  68027.75183387,  44136.37698061,
  44020.27071045,  51727.60024216,  82542.47600601,  42100.01220024]
```

```
In [19]: from sklearn.metrics import r2_score
r2_score(y_test,ypred)
```

Out[19]: 0.7691178949173114

```
In [20]: from sklearn.metrics import mean_squared_error
n=mean_squared_error(ypred,y_test)
```

```
In [21]: #Results=pd.DataFrame(columns=['Actual','Predicted'])
#Results['Actual']=y_test
Results=pd.DataFrame(columns=['Price','Predicted'])
Results['Price']=y_test
Results['Predicted']=ypred
#Result['km']=x_test['km']
Results=Results.reset_index()
Results['Id']=Results.index
Results.head(15)
```


Out[21]:

	index	Price	Predicted	Id
0	481	120000	75731.244523	0
1	76	103000	66029.847723	1
2	1502	32473	34401.236097	2
3	669	29000	47789.215452	3
4	1409	18800	31445.263386	4
5	1414	39751	28940.452249	5
6	1089	33160	29985.553110	6
7	1507	17324	27782.989861	7
8	970	29000	18389.948101	8
9	1198	38000	43647.331735	9
10	1088	11511	26400.685613	10
11	576	78000	65315.517185	11
12	965	50247	78520.353231	12
13	1488	108000	79990.405926	13
14	1432	42095	40933.248327	14

```
In [22]: import seaborn as sns  
import matplotlib.pyplot as plt
```

```
In [23]: sns.lineplot(x='Id',y='Price',data=Results.head(50))#blue  
sns.lineplot(x='Id',y='Predicted',data=Results.head(50))#orange  
plt.plot()
```

Out[23]: []



```
In [24]: Results['DIFF']=Results.apply(lambda row: row.Price-row.Predicted,axis=1)
```


In [25]: Results

Out[25]:

	index	Price	Predicted	Id	DIFF
0	481	120000	75731.244523	0	44268.755477
1	76	103000	66029.847723	1	36970.152277
2	1502	32473	34401.236097	2	-1928.236097
3	669	29000	47789.215452	3	-18789.215452
4	1409	18800	31445.263386	4	-12645.263386
...
503	291	22000	15885.136964	503	6114.863036
504	596	85500	100920.913776	504	-15420.913776
505	1489	22148	27354.930930	505	-5206.930930
506	1436	61000	74399.449201	506	-13399.449201
507	575	19112	12967.563730	507	6144.436270

508 rows × 5 columns

ridge regression

```
In [26]: from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import Ridge
alpha=[1e-15,1e-10,1e-8,1e-4,1e-3,1e-2,1,5,10,20,30]
ridge=Ridge()
parameters={'alpha':alpha}
ridge_regressor=GridSearchCV(ridge,parameters)
ridge_regressor.fit(x_train,y_train)
```

```
Out[26]: GridSearchCV(estimator=Ridge(),
                      param_grid={'alpha': [1e-15, 1e-10, 1e-08, 0.0001, 0.001, 0.01, 1,
                                             5, 10, 20, 30]})
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [27]: ridge_regressor.best_params_
```

```
Out[27]: {'alpha': 30}
```

```
In [28]: ridge=Ridge(alpha=30)
ridge.fit(x_train,y_train)
y_pred_ridge=ridge.predict(x_test)
```

```
In [29]: from sklearn.metrics import mean_squared_error
Ridge_Error=mean_squared_error(y_pred_ridge,y_test)
Ridge_Error
```

```
Out[29]: 346741563.1891953
```

```
In [30]: from sklearn.metrics import r2_score
r2_score(y_test,y_pred_ridge)
```

```
Out[30]: 0.7691472836651715
```

```
In [31]: Results=pd.DataFrame(columns=['Actual','predicted'])
Results['Actual']=y_test
Results['predicted']=y_pred_ridge
#Result['km']=x_test['km']
Results=Results.reset_index()
Results['Id']=Results.index
Results.head(10)
```


```
Out[31]:
```

	index	Actual	predicted	Id
0	481	120000	75610.216839	0
1	76	103000	66438.222154	1
2	1502	32473	34390.439217	2
3	669	29000	47783.153021	3
4	1409	18800	31444.515496	4
5	1414	39751	28942.281540	5
6	1089	33160	29995.447469	6
7	1507	17324	27781.363705	7
8	970	29000	18397.986372	8
9	1198	38000	43663.963925	9

```
In [32]: import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [33]: sns.lineplot(x='Id',y='Actual',data=Results.head(50))#blue  
sns.lineplot(x='Id',y='predicted',data=Results.head(50))#orange  
plt.plot()
```

Out[33]: []



elastic net

```
In [34]: from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import ElasticNet

elastic = ElasticNet()

parameters = {'alpha': [1e-15, 1e-10, 1e-8, 1e-4, 1e-3, 1e-2, 1, 5, 10, 20]}

elastic_regressor=GridSearchCV(elastic, parameters)

elastic_regressor.fit(x_train, y_train)
```

```
Out[34]: GridSearchCV(estimator=ElasticNet(),
                      param_grid={'alpha': [1e-15, 1e-10, 1e-08, 0.0001, 0.001, 0.01, 1,
                                             5, 10, 20]})
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [35]: elastic_regressor.best_params_
```

```
Out[35]: {'alpha': 1}
```

```
In [36]: elastic=ElasticNet(alpha=0.01)
elastic.fit(x_train,y_train)
y_pred_elastic=elastic.predict(x_test)
```

```
In [37]: from sklearn.metrics import r2_score
r2_score(y_test,y_pred_elastic)
```

```
Out[37]: 0.7691157601736113
```

```
In [38]: from sklearn.metrics import mean_squared_error
elastic_Error=mean_squared_error(y_pred_elastic,y_test)
elastic_Error
```

Out[38]: 346788911.580474

```
In [39]: Results=pd.DataFrame(columns=['Actual','predicted'])
Results['Actual']=y_test
Results['predicted']=y_pred_elastic
#Result['km']=x_test['km']
Results=Results.reset_index()
Results['price']=Results.index
Results.head(10)
```


Out[39]:

	index	Actual	predicted	price
0	481	120000	75710.809218	0
1	76	103000	66113.691091	1
2	1502	32473	34399.544383	2
3	669	29000	47789.104506	3
4	1409	18800	31445.472680	4
5	1414	39751	28941.068113	5
6	1089	33160	29987.853086	6
7	1507	17324	27782.851276	7
8	970	29000	18391.334150	8
9	1198	38000	43651.631475	9

```
In [40]: import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [41]: sns.lineplot(x='price',y='Actual',data=Results.head(50))#blue  
sns.lineplot(x='price',y='predicted',data=Results.head(50))#orange  
plt.plot()
```

Out[41]: []



In []:

